

## 34. Imbalanced Dataset

- Imbalanced dataset means that your data consists of multi categories and one category is repetitive in the data
- model is biased due to repetition of one category in the data
- Suppose your data consist of 500 rows:
  - 400 rows for cat and 100 rows for dog,
  - so the model will be biased towards cat

### 34.1 Techniques to handle imbalanced data

#### 34.1.1 Random Under Sampling

- we will reduce the majority of the class so that it will have same number of as the minority
- for example out of 500 rows for cats and dogs, we will reduce (randomly) the rows to 100 for cats that is equal to 100 rows of dogs

#### 34.1.2 Random Over Sampling

- We will increase the size of manority is inactive class to the size of majority calss i.e. active
- for example out of 500 rows for cats and dogs, we will repeat/duplicate (randomly) the rows for dogs to make it to 400 that is equal to 400 rows of cats

```
In [1]: import pandas as pd
```

```
In [3]: dataset = pd.read_csv(r'Data/Social_Network_Ads.csv')  
dataset.head(3)
```

```
Out[3]:
```

	Age	EstimatedSalary	Purchased
0	19	19000	0
1	35	20000	0
2	26	43000	0

**To check the data if it is imbalanced or not**

```
In [4]: dataset['Purchased'].value_counts()
```

```
Out[4]: 0    257
        1    143
        Name: Purchased, dtype: int64
```

So hence the data is **imbalanced** b/c both categories are not equal, so the data will be biased towards 0

```
In [12]: x = dataset.iloc[:, :-1]
        x
```

```
Out[12]:
```

	Age	EstimatedSalary
0	19	19000
1	35	20000
2	26	43000
3	27	57000
4	19	76000
...	...	...
395	46	41000
396	51	23000
397	50	20000
398	36	33000
399	49	36000

400 rows × 2 columns

```
In [13]: y = dataset['Purchased']
        y
```

```
Out[13]: 0      0
        1      0
        2      0
        3      0
        4      0
        ..
        395    1
        396    1
        397    1
        398    0
        399    1
        Name: Purchased, Length: 400, dtype: int64
```

```
In [14]: from sklearn.model_selection import train_test_split
```

```
In [15]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.20, random_st
```

```
In [16]: from sklearn.linear_model import LogisticRegression
```

```
In [17]: lg = LogisticRegression()  
lg.fit(x_train, y_train)
```

```
Out[17]: ▾ LogisticRegression  
LogisticRegression()
```

```
In [22]: lg.score(x_test, y_test)*100
```

```
Out[22]: 65.0
```

```
In [23]: # y_true is 0  
lg.predict([[19, 19000]])
```

C:\Users\rashi\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names  
warnings.warn(

```
Out[23]: array([0], dtype=int64)
```

```
In [25]: # y_true is 1  
lg.predict([[45, 26000]])
```

C:\Users\rashi\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names  
warnings.warn(

```
Out[25]: array([0], dtype=int64)
```

It has given wrong prediction, Reason: B/c the input data is **imbalanced**

```
In [26]: # y_true is 1  
lg.predict([[46, 28000]])
```

C:\Users\rashi\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names  
warnings.warn(

```
Out[26]: array([0], dtype=int64)
```

Again, it has given wrong prediction, Reason: B/c the input data is **imbalanced**

```
In [ ]:
```

**So hence we will balance our data by either:**

- random under sampling, or
- random over sampling

## 34.2.1 Balacing the data by Random Under Sampling (Practical)

```
In [28]: from imblearn.under_sampling import RandomUnderSampler
```

```
In [31]: ru = RandomUnderSampler()  
ru_x, ru_y = ru.fit_resample(x,y)
```

```
In [32]: ru_x
```

```
Out[32]:
```

	Age	EstimatedSalary
<b>224</b>	35	60000
<b>49</b>	31	89000
<b>153</b>	36	50000
<b>132</b>	30	87000
<b>359</b>	42	54000
...	...	...
<b>393</b>	60	42000
<b>395</b>	46	41000
<b>396</b>	51	23000
<b>397</b>	50	20000
<b>399</b>	49	36000

286 rows × 2 columns

```
In [33]: ru_y
```

```
Out[33]:
```

224	0
49	0
153	0
132	0
359	0
..	
393	1
395	1
396	1
397	1
399	1

Name: Purchased, Length: 286, dtype: int64

**Now after applying under sampling technique we will see, if 0 count is reduced to 143 or not**

```
In [34]: # Remember, our original data has following counts:
dataset['Purchased'].value_counts()
```

```
Out[34]: 0    257
         1    143
         Name: Purchased, dtype: int64
```

```
In [35]: ru_y.value_counts()
```

```
Out[35]: 0    143
         1    143
         Name: Purchased, dtype: int64
```

**So you can see 0 has reduced to 143 and now our data is balanced!**

Now we have new data variables that are **ru\_x, ru\_y**

**We will apply logistic regression on this new dataset that is balanced data**, so we first split the data into train and test and then will apply logistic regression model

```
In [36]: from sklearn.model_selection import train_test_split
```

```
In [37]: x_train, x_test, y_train, y_test = train_test_split(ru_x, ru_y, test_size=0.20, ran
```

```
In [38]: from sklearn.linear_model import LogisticRegression
```

```
In [39]: ru_lg = LogisticRegression()
ru_lg.fit(x_train, y_train)
```

```
Out[39]: ▾ LogisticRegression
LogisticRegression()
```

```
ru_lg.score(x_test, y_test)*100
```

**To check if the model has improved or not** we will supply same value as were predicted wrongly

```
In [41]: # y_true is 1
ru_lg.predict([[45, 26000]])
```

```
C:\Users\rashi\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names
warnings.warn(
```

```
Out[41]: array([1], dtype=int64)
```

**Hurrahh, now it has given accurate prediction**, lets try second test..

```
In [43]: # y_true is 1
ru_lg.predict([[46, 28000]])
```

```
C:\Users\rashi\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names
warnings.warn(
```

```
Out[43]: array([1], dtype=int64)
```

**Oh yes, the second prediction is also accurate!!!**

```
In [44]: # y_true is 0
lg.predict([[19, 19000]])
```

```
C:\Users\rashi\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names
warnings.warn(
```

```
Out[44]: array([0], dtype=int64)
```

**Great, accurate prediction again!!!**

**Conclusion is our model is not more biased**

## 34.2.2 Balacing the data by Random Over Sampling (Practical)

```
In [46]: from imblearn.over_sampling import RandomOverSampler
```

```
In [47]: ro = RandomOverSampler()
ro_x, ro_y = ro.fit_resample(x,y)
```

```
In [48]: ro_x
```

Out[48]:

	Age	EstimatedSalary
0	19	19000
1	35	20000
2	26	43000
3	27	57000
4	19	76000
...	...	...
509	42	73000
510	55	39000
511	46	28000
512	37	93000
513	46	79000

514 rows × 2 columns

In [49]:

```
ro_y
```

Out[49]:

```
0      0
1      0
2      0
3      0
4      0
..
509    1
510    1
511    1
512    1
513    1
```

Name: Purchased, Length: 514, dtype: int64

In [61]:

```
# Remember, our original data has following counts:
dataset['Purchased'].value_counts()
```

Out[61]:

```
0      257
1      143
```

Name: Purchased, dtype: int64

So 1 should be increased to 257 as well as we have applied random over sampling method

In [60]:

```
ro_y.value_counts()
```

Out[60]:

```
0      257
1      257
```

Name: Purchased, dtype: int64

Now input is ro\_x and output is ro\_y, we will split the data into test and train and then apply logistic regression model

```
In [50]: x_train, x_test, y_train, y_test = train_test_split(ro_x, ro_y, test_size=0.20, ran
```

```
In [52]: ro_lg = LogisticRegression()  
ro_lg.fit(x_train, y_train)
```

```
Out[52]: ▾ LogisticRegression  
LogisticRegression()
```

```
In [54]: ro_lg.score(x_test, y_test)*100
```

```
Out[54]: 88.3495145631068
```

**Accuracy of the model is increased, impressive!!**

```
In [56]: # y_true is 0  
lg.predict([[19, 19000]])
```

```
C:\Users\rashi\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names  
warnings.warn(
```

```
Out[56]: array([0], dtype=int64)
```

```
In [57]: # y_true is 1  
ru_lg.predict([[46, 28000]])
```

```
C:\Users\rashi\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names  
warnings.warn(
```

```
Out[57]: array([1], dtype=int64)
```

```
In [58]: # y_true is 1  
ru_lg.predict([[45, 26000]])
```

```
C:\Users\rashi\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names  
warnings.warn(
```

```
Out[58]: array([1], dtype=int64)
```

**Conclusion:**

- All predictions are accurate by even Random Over Sampling
- in case of Random Over Sampling, the model accuracy is also increased from 65% (on imbalanced data) to 88% (on balanced data)



- in case of Random Under Sampling, the model accuracy is also decreased from 65% (on imbalanced data) to 58% (on balanced data)
- However, the model is predicting accurately after making the data balanced by both methods, i.e, Random over sampling, Random under sampling

In [ ]: